

SDIC Lab 4: MapReduce

Introduction

In this lab you'll build a MapReduce library as an introduction to programming in Java and to building fault tolerant distributed systems. In the first part you will write a simple MapReduce program. In the second part you will write a Master that hands out tasks to MapReduce workers. And you will complete a parallel version of Mapreduce and handle the failures of workers in part 3 and part 4. There is a optional exercise in part 5, which would you bonus if you complete a inverted index application.

The interface to the library and the approach to fault tolerance is similar to the one described in the original [MapReduce paper](#).

Copyright

This is a Java-implementation of MIT 6.824 Lab 1: MapReduce, for SDIC course lab, SJTU. Some resources directly or indirectly refer to the [original materials](#).

Notice: For educational or personal usage only.

Collaboration Policy

You must write all the code you hand in for this course, except for code that we give you as part of the assignment. You are not allowed to look at anyone else's solution, which means that you may discuss the assignments with other students, but you may not look at or copy each others' code. The reason for this rule is that we believe you will learn the most by designing and implementing your lab solution code yourself.

Software

You'll implement this lab (and all the labs) in **Java**. Java is an mature OOP language and many learning materials could be found online. We believe that most of you are able to learn it and work out this lab with it. We will grade your labs using [Java 1.8 latest](#), running on Java HostSpot VM(64 bit). You better use this too, though we don't know if there is any problem with other versions.

Java is platform-independent so you don't need to worry about the difference between OS and hardwares. But it needs to ~~set environment variables~~. Google it and learn how to do it for your OS.

We supply you with parts of a MapReduce implementation that supports both distributed and non-distributed operation (just the boring bits). You'll fetch the initial lab software from [here](#). This is a Maven project under IDEA. To learn more about Maven, look at the [Apache Maven](#) or other resources. [IDEA](#) is the most widely used Java IDE now, we recommend you to use this IDE because it accelerates the developing, though it is not required, some following steps are demonstrated by IDEA.

The Map/Reduce implementation we give you has support for two modes of operation, *sequential* and *distributed*. In the former, the map and reduce tasks are executed one at a time: first, the first map task is executed to completion, then the second, then the third, etc. When all the map tasks have finished, the first reduce task is run, then the second, etc. This mode, while not very fast, is useful for debugging. The distributed mode runs many worker threads that first execute map tasks in parallel, and then reduce tasks. This is much faster, but also harder to implement and debug.

Preamble: Getting familiar with the source

The mapreduce project provides a simple Map/Reduce library (in the common, core, rpc package). Applications should normally call *distributed()* [located in core/Master.java] to start a job, but may instead call *sequential()* [also in core/Master.java] to get a sequential execution for debugging.

The code executes a job as follows:

1. The application provides a number of input files, a map function, a reduce function, and the number of reduce tasks (`nReduce`).
2. A master is created with this knowledge. It starts an RPC server, and waits for workers to register (using the RPC call `register()` [defined in `core/Master.java`]). As tasks become available (in steps 4 and 5), `schedule()` [`core/Scheduler.java`] decides how to assign those tasks to workers, and how to handle worker failures.
3. The master considers each input file to be one map task, and calls `doMap()` [`core/Mapper.java`] at least once for each map task. It does so either directly (when using `sequential()`) or by issuing the `doTask` RPC to a worker [`core/Worker.java`]. Each call to `doMap()` reads the appropriate file, calls the map function on that file's contents, and writes the resulting key/value pairs to `nReduce` intermediate files. `doMap()` hashes each key to pick the intermediate file and thus the reduce task that will process the key. There will be `nMap` x `nReduce` files after all map tasks are done. Each file name contains a prefix, the map task number, and the reduce task number. If

there are two map tasks and three reduce tasks, the map tasks will create these six intermediate files:

```
1 mrtmp.xxx-0-0
2 mrtmp.xxx-0-1
3 mrtmp.xxx-0-2
4 mrtmp.xxx-1-0
5 mrtmp.xxx-1-1
6 mrtmp.xxx-1-2
```

Each worker must be able to read files written by any other worker, as well as the input files. Real deployments use distributed storage systems such as GFS to allow this access even though workers run on different machines. In this lab you'll run all the workers on the same machine, and use the local file system.

4. The master next calls `doReduce()` [`core/Reducer.java`] at least once for each reduce task. As with `doMap()`, it does so either directly or through a worker. The `doReduce()` for reduce task `r` collects the `r`'th intermediate file from each map task, and calls the reduce function for each key that appears in those files. The reduce tasks produce `nReduce` result files.
5. The master calls `merge()` [`core/Master.java`], which merges all the `nReduce` files produced by the previous step into a single output.
6. The master sends a Shutdown RPC to each of its workers, and then shuts down its own RPC server.

Note: Over the course of the following exercises, you will have to write/modify `doMap`, `doReduce`, and `schedule` yourself. These are located in `core/Mapper.java`, `core/Reducer.java`, and `core/Scheduler.java` respectively. You will also have to write the map and reduce functions in `WordCount.java`.

You should not modify any other files, but reading them might be useful in order to understand how the other methods fit into the overall architecture of the system.

Part I: Map/Reduce input and output

The Map/Reduce implementation you are given is missing some pieces. Before you can write your first Map/Reduce function pair, you will need to fix the sequential implementation. In particular, the code we give you is missing two crucial pieces: the function that divides up the output of a map task, and the function that gathers all the inputs for a reduce task. These tasks are carried out by the `doMap()` function in `Mapper.java`, and the `doReduce()` function in `Reducer.java` respectively. The comments in those files should point you in the right direction.

To help you determine if you have correctly implemented `doMap()` and `doReduce()`, we have provided you with a test suite that checks the correctness of your implementation. These tests are implemented in the file `MRTTest.java`. To run the tests for the sequential implementation that you have now fixed, run:

```
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

TASK: You receive full credit for this part if your software passes the Sequential tests (as run by the command above) when we run your software on our machines.

If the output did not show `ok` next to the tests, your implementation has a bug in it. To give more verbose output, set `debugEnabled = true` in `common/Utils.java`. You will get much more output along the lines of:

Tests passed: 2 of 2 tests – 5 s 324 ms		
✓ MRTTest	5 s 324 ms	master: Starting Map/Reduce task test
✓ testSequentialSingle	3 s 529 ms	Merge: read mrtmp.test-res-0
✓ testSequentialMany	1 s 795 ms	master: Map/Reduce task completed
✓ Tests passed: 2 of 2 tests – 5 s 324 ms		
✓ MRTTest	5 s 324 ms	master: Starting Map/Reduce task test
✓ testSequentialSingle	3 s 529 ms	Merge: read mrtmp.test-res-0
✓ testSequentialMany	1 s 795 ms	Merge: read mrtmp.test-res-1
		Merge: read mrtmp.test-res-2
		master: Map/Reduce task completed

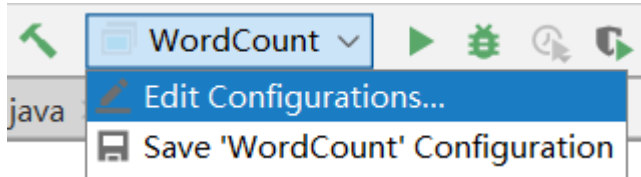
Part II: Single-worker word count

Now you will implement word count — a simple Map/Reduce example. Look in `WordCount.java`; you'll find empty `mapFunc()` and `reduceFunc()` functions. Your job is to insert code so that `WordCount.java` reports the number of occurrences of each word in its input. A word is any contiguous sequence of letters, as determined by regular expression:

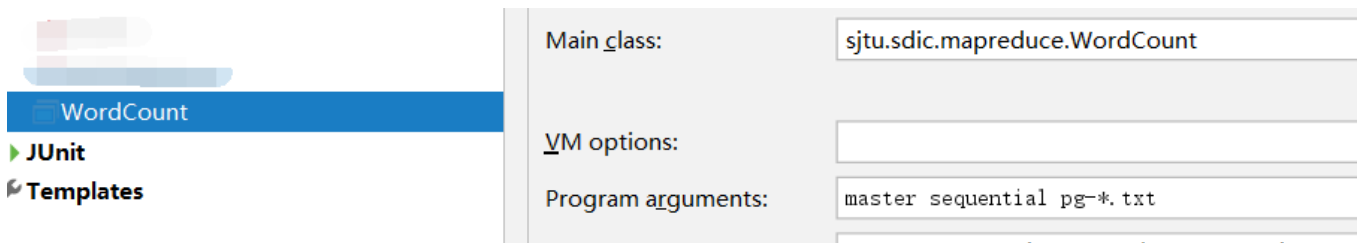
```
1 [a-zA-Z0-9]+
```

There are some input files with pathnames of the form `pg-*.txt` in project's root dir, downloaded from [Project Gutenberg](#). Here's how to run `WordCount.java` with the input files:

1. Find "Edit Configurations" of WordCount



2. Edit Program arguments as below. This is same with specifying parameters in Console/shell.



Review Section 2 of the [MapReduce paper](#). Your `mapFunc()` and `reduceFunc()` functions will differ a bit from those in the paper's Section 2.1. Your `mapFunc()` will be passed the name of a file, as well as that file's contents; it should split the contents into words, and return a List of `common/KeyValue`. While you can choose what to put in the keys and values for the `mapFunc` output, for word count it only makes sense to use words as the keys. Your `reduceFunc()` will be called once for each key, with a slice of all the values generated by `mapFunc()` for that key. It must return a string containing the total number of occurrences of the key.

- **Hint:** Use `Pattern` and `Matcher` to extract words.
- **Hint:** `Integer.valueOf(String str)` could parse a valid string to integer.

The output will be in the file "mrtmp.wcseq". Your implementation is correct if the following command produces the output shown here (under XNIX):

```
$ sort -n -k2 mrtmp.wcseq | tail -10
that: 7871
it: 7987
in: 8415
was: 8578
a: 13382
of: 13536
I: 14296
to: 16079
and: 23612
the: 29748
```

TASK: You receive full credit for this part if your Map/Reduce word count output matches the correct output for the sequential execution above when we run your software on our machines.

Part III: Distributing MapReduce tasks

Your current implementation runs the map and reduce tasks one at a time. One of Map/Reduce's biggest selling points is that it can automatically parallelize ordinary sequential code without any extra work by the developer. In this part of the lab, you will complete a version of MapReduce that splits the work over a set of worker threads that run in parallel on multiple cores. While not distributed across multiple machines as in real Map/Reduce deployments, your implementation will use RPC to simulate distributed computation.

The code in `core/Master.java` does most of the work of managing a MapReduce job. We also supply you with the complete code for a worker thread, in `core/Worker.java`, as well as some code to deal with RPC in package `rpc`.

Your job is to implement `schedule()` in `core/Scheduler.java`. The master calls `schedule()` twice during a MapReduce job, once for the Map phase, and once for the Reduce phase. `schedule()`'s job is to hand out tasks to the available workers. There will usually be more tasks than worker threads, so `schedule()` must give each worker a sequence of tasks, one at a time. `schedule()` should wait until all tasks have completed, and then return. `schedule()` learns about the set of workers by reading its `registerChan` argument. That channel yields a string for each worker, containing the worker's RPC address. Some workers may exist before `schedule()` is called, and some may start while `schedule()` is running; all will appear on `registerChan`. `schedule()` should use all the workers, including ones that appear after it starts.

`schedule()` tells a worker to execute a task by calling

```
1 Call.getWorkerRpcService(worker).doTask(arg);
```

This RPC's arguments are defined by `DoTaskArgs` in `common/DoTaskArgs.java`. The `file` element is only used by Map tasks, and is the name of the file to read; `schedule()` can find these file names in `mapFiles`.

Your solution to Part III should only involve modifications to `Scheduler.java`. If you modify other files as part of debugging, please restore their original contents and then test before submitting.

The corresponding test methods are as below, defined in `MRTest.java`:

```
@Test
public void testParallelBasic() {

@Test
public void testParallelCheck() {
```

The latter verifies that your scheduler causes workers to execute tasks in parallel.

TASK: You will receive full credit for this part if your software passes `testParallelBasic` and `testParallelCheck` when we run your software on our machines.

- **Hint:** RPC implementation relies on `SO FARPC`, documents could be found [here](#).
- **Hint:** `schedule()` should send RPCs to the workers in parallel so that the workers can work on tasks concurrently. Learn how to use Java Thread API.
- **Hint:** `schedule()` must wait for a worker to finish before it can give it another task. Try to exploit `Channel`.
- **Hint:** You may find `CountDownLatch` is useful.
- **Hint:** To stop a thread, using `interrupt()`. Google it to learn more.
- **Hint:** The easiest way to track down bugs is to insert print statements (perhaps calling `debug()` in `Utils.java`, and then think about whether the output matches your understanding of how your code should behave.

Part IV: Handling worker failures


In this part you will make the master handle failed workers. MapReduce makes this relatively easy because workers don't have persistent state. If a worker fails while handling an RPC from the master, the master's `call()` will eventually return `false` due to a timeout. In that situation, the master should re-assign the task given to the failed worker to another

worker.

An RPC failure doesn't necessarily mean that the worker didn't execute the task; the worker may have executed it but the reply was lost, or the worker may still be executing but the master's RPC timed out. Thus, it may happen that two workers receive the same task, compute it, and generate output. Two invocations of a map or reduce function are required to generate the same output for a given input (i.e. the map and reduce functions are "functional"), so there won't be inconsistencies if subsequent processing sometimes reads one output and sometimes the other. In addition, the MapReduce framework ensures that map and reduce function output appears atomically: the output file will either not exist, or will contain the entire output of a single execution of the map or reduce function (the lab code doesn't actually implement this, but instead only fails workers at the end of a task, so there aren't concurrent executions of a task).

Note: You don't have to handle failures of the master. Making the master fault-tolerant is more difficult because it keeps state that would have to be recovered in order to resume operations after a master failure. Much of the later labs are devoted to this challenge.

Your implementation must pass the two remaining test cases in `MRTest.java`. The first case tests the failure of one worker, while the second test case tests handling of many failures of workers. Periodically, the test cases start new workers that the master can use to make forward progress, but these workers fail after handling a few tasks. To run these tests:



```
@Test
public void testOneFailure() {
    ..
}

@Test
public void testManyFailures() {
```

TASK: You receive full credit for this part if your software passes the tests with worker failures (those run by the command above) when we run your software on our machines.

Your solution to Part IV should only involve modifications to `Scheduler.java`. If you modify other files as part of debugging, please restore their original contents and then test before submitting.

Part V: Inverted index generation (optional, as bonus)

For this optional no-credit exercise, you will build Map and Reduce functions for generating an *inverted index*.

Inverted indexes are widely used in computer science, and are particularly useful in document searching. Broadly speaking, an inverted index is a map from interesting facts about the underlying data, to the original location of that data. For example, in the context of search, it might be a map from keywords to documents that contain those words.

We have created a second binary in `InvertedIndex.java` that is very similar to the `WordCount.java` you built earlier. You should modify `mapFunc` and `reduceFunc` in `InvertedIndex.java` so that they together produce an inverted index. Running `InvertedIndex.java` should output a list of tuples, one per line, in the following format:

```
1 $ head -n5 mrtmp.iiseg
2 A: 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg
3 ABOUT: 1 pg-tom_sawyer.txt
4 ACT: 1 pg-being_ernest.txt
5 ACTRESS: 1 pg-dorian_gray.txt
6 ACTUAL: 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.t
```

If it is not clear from the listing above, the format is:

```
1 word: #documents documents,sorted,and,separated,by,commas
```

The expected output should be like this (under XNIX):

```

1 ALL=C sort -k1,1 mrtmp.iiseq sort -snk2,2 grep -v '16' tail -10
2 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gri
3 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gr
4 : 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-c
5 rday: 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,
6 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gri
7 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gri
8 : 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-c
9 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gr
10 elf: 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,r
11 8 pg-being_ernest.txt,pg-dorian_gray.txt,pg-frankenstein.txt,pg-gri

```

Handin procedure

Important:

Before submitting, please make sure all test methods are passed in EVERY attempts.

Submit your:

Mapper.java, Reducer.java, Scheduler.java, WordCount.java, InvertedIndex.java(if you have done it) to linfusheng19@163.com.

The requirement format for naming the submitted file : **SDICLab4_StudentID_Name.rar**

DEADLINE: 23:59 of June 6, 2019

Note: You may submit multiple times. We will use the timestamp of your lastsubmission for the purpose of calculating late days.